



Delta NLA RL training curves — full-transformer delta

3,000 RL steps on one H100 · raw per-step traces beneath trailing 50-step means · held-out evaluation every 250 steps

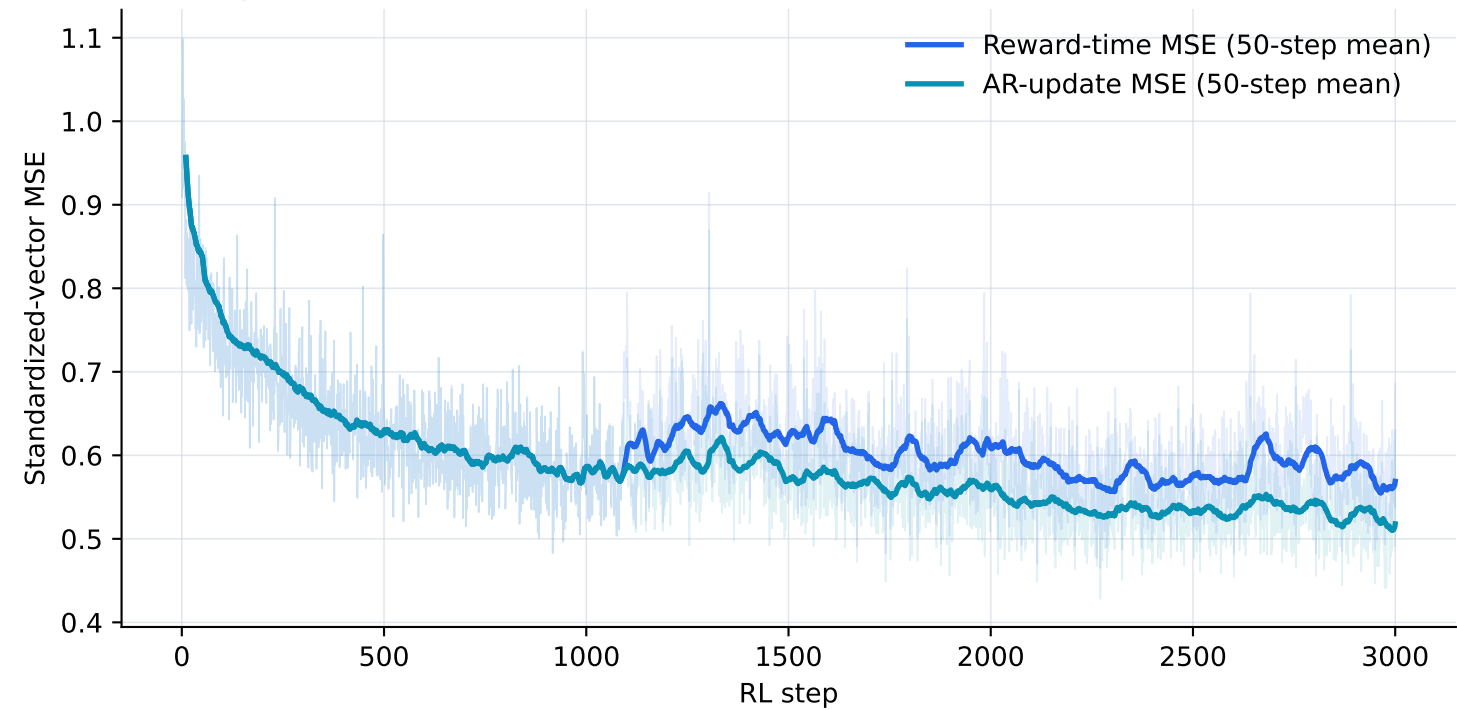
Held-out reconstruction improves 0.121 → 0.501



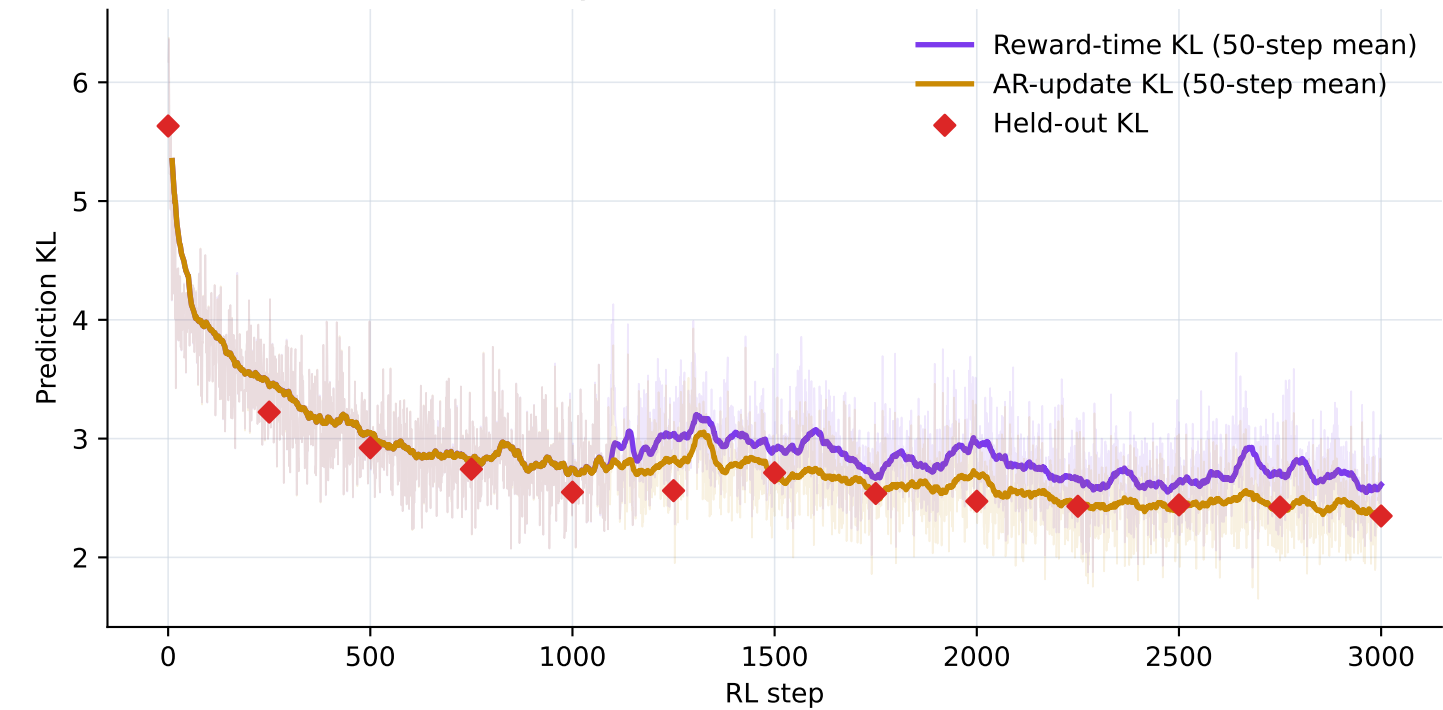
Held-out objective (lower is better)



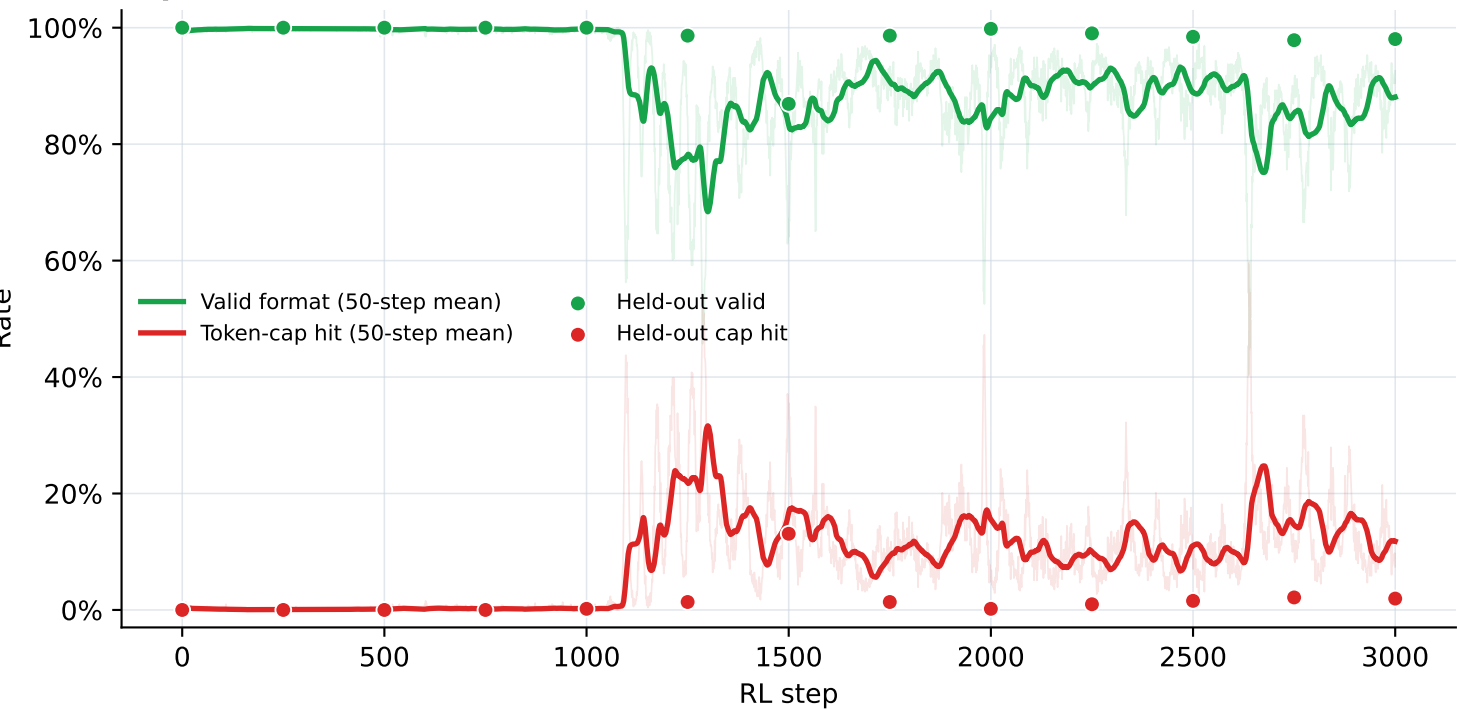
Training reconstruction error



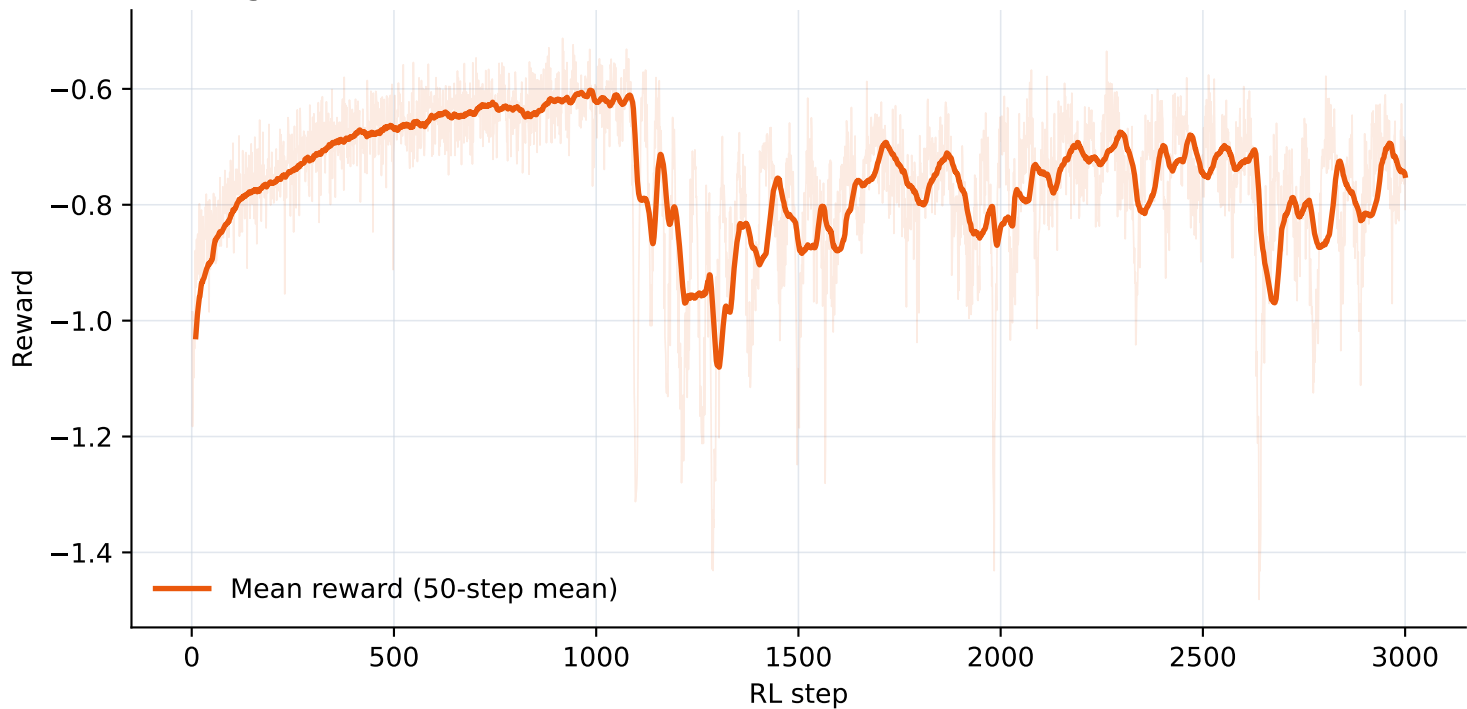
Next-token prediction fidelity



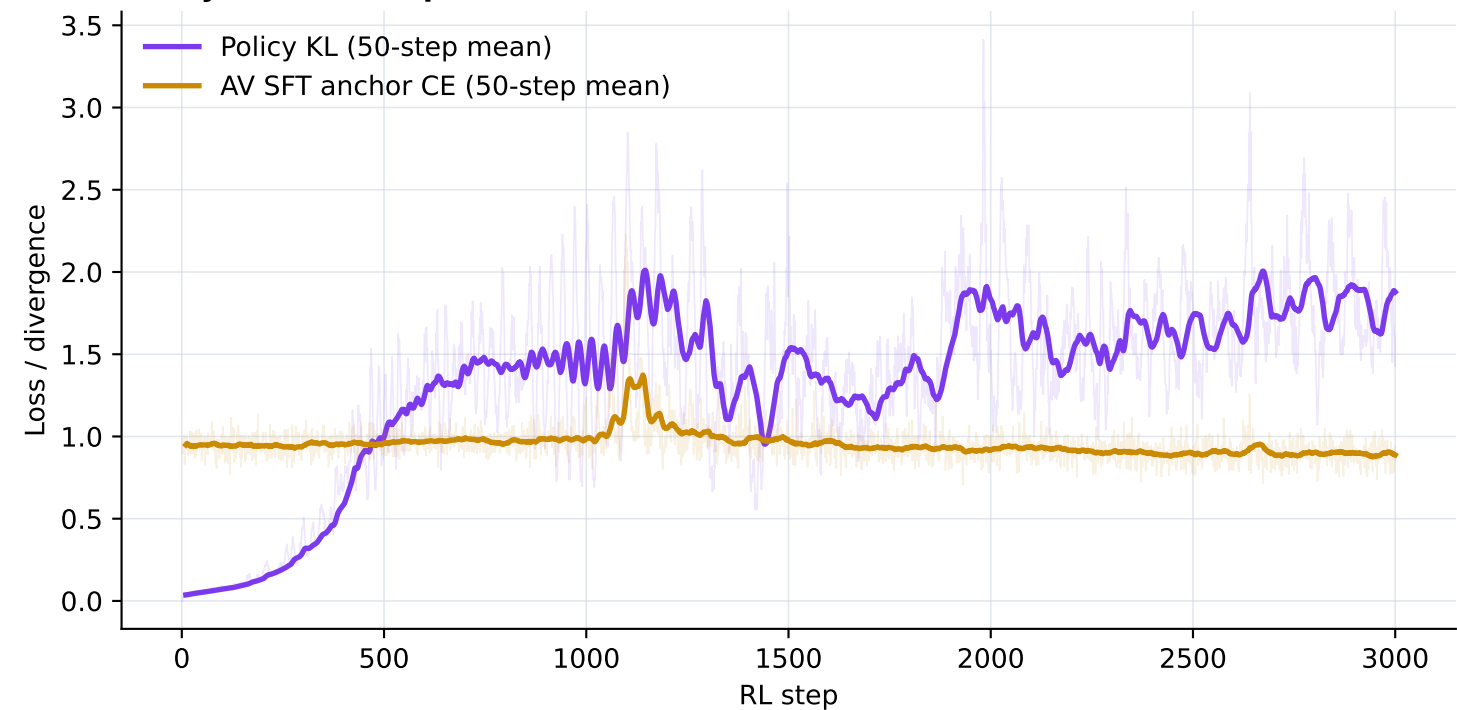
Explanation format and truncation



Training reward



Policy drift and supervised anchor



Generation length and AR training coverage

